


# Computer Vision

Modelisation

# Reminder: purpose

**Computer vision** is a field of artificial intelligence that trains computers to interpret and understand the visual world. Using digital images from cameras and videos and **deep learning** models, machines can accurately **identify and classify objects** — and then **react** to what they “see.”

- Object detection
- Object tracking
- Image segmentation / search
- 3D scene modeling



# Reminder: problems




- Object orientation
  - Lighting conditions
  - Obstruction
  - ...
- ➔ **Too many possibilities !**




# Why deep learning ?

## Machine Learning



## Deep Learning



# Plan

## Classification Results (CLS)



➤ **Tasks**

➤ **CNN**

➤ **Compilation**

➤ **Transfert learning**

# Tasks


Classification - detection - segmentation - creation



CAT



CAT



CAT, DOG, DUCK



CAT, DOG, DUCK

Single object

Multiple objects

# Classification

## Particularity:

- N possible outputs with  $N = \text{number of classes}$ .



CAT

## Algorithms

- Mark 1
- SVM
- CNN

## Datasets

- ImageNet
- CIFAR 100
- ISIC
- MURA
- DermNet


## Models

- LeNet5
- VGG
- GoogleNet / MobileNet
- ResNet

# Object detection

## Particularity:

- Is an object present ?
- Location:
  - x and y (top left of the object area, or center)
  - Area height and width



CAT

## Algorithms

- Sliding window
- CNN

## Datasets

- ImageNet
- COCO
- PASCAL

## Models

- R-CNN
- YOLO
- SSD



# Image segmentation

## Particularity:

- Is an object present ?
- Object boundaries
- Give a value for each pixel



CAT, DOG, DUCK

## Algorithms

- CNN

## Datasets

- COCO
- PASCAL
- BraTS
- Agriculture-Vision
- SpaceNet

## Models

- U-net
- Mask R-CNN

# Image creation / generation

## Particularity:

- Give a value for each pixel
- Unlike any training image
- Improve resolution

This Person Does Not Exist

This Person Does Not Exist

This Person Does Not Existthispersondoesnotexist.com



## Algorithms

- CNN

## Datasets

- Any...

## Models

- GANs
- VAEs

# Potential

- Biomédical
- Robotics
- Airports
- Security / Surveillance
- Space exploration
- Customer experience
- ...


## Behavioral tracking



# Potential

- Biomédical
- Robotics
- Airports
- Security / Surveillance
- Space exploration
- Customer experience
- ...


## Healthcare



# Potential

## Autonomous vehicles

- Biomédical
- Robotics
- Airports
- Security / Surveillance
- Space exploration
- Customer experience
- ...




# CNN

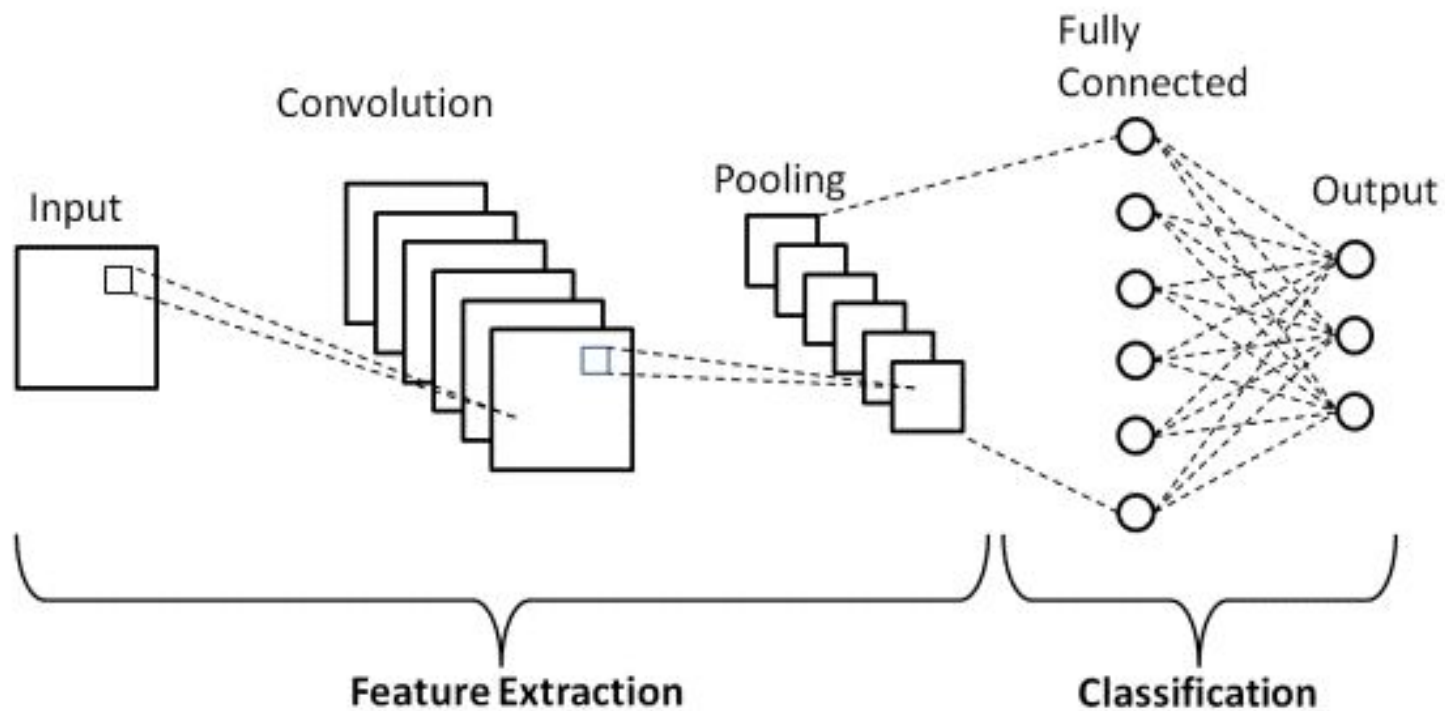
“Learn” the features

# MLP

Feature detectors			Feature descriptors		
250	249	249	248	242	
246	<b>A</b>	44	243	(240, 230, ...)	
244	<b>B</b>	42	240	(190, 210, ...)	
241	<b>C</b>	40	240	(40, 150, ...)	
...	<b>D</b>	..	...	(200, 180, ...)	





# CNN





# CNN: convolution layer



## Activation function:

- ReLU:  $f(x) = \text{Max}(0, x)$

# CNN: convolution layer

## Image (input):

- W : width (pixels)
- H : height (pixels)
- D : number of channels

## Hyperparameters:

- K : number of filters
- F : filter size in pixels (F x F x D)
- S : sliding step (pixels)
- P : 0 padding (pixels)
- Initialization: random / He / Gorot ...


## Feature maps (output):

- $W_c : (W - F + 2P) / S + 1$
- $H_c : (H - F + 2P) / S + 1$
- $D_c : K$

## Common choice:

- F = 3 or 5
- S = 1 or 2
- P = 1

# CNN



# CNN: Batch normalization


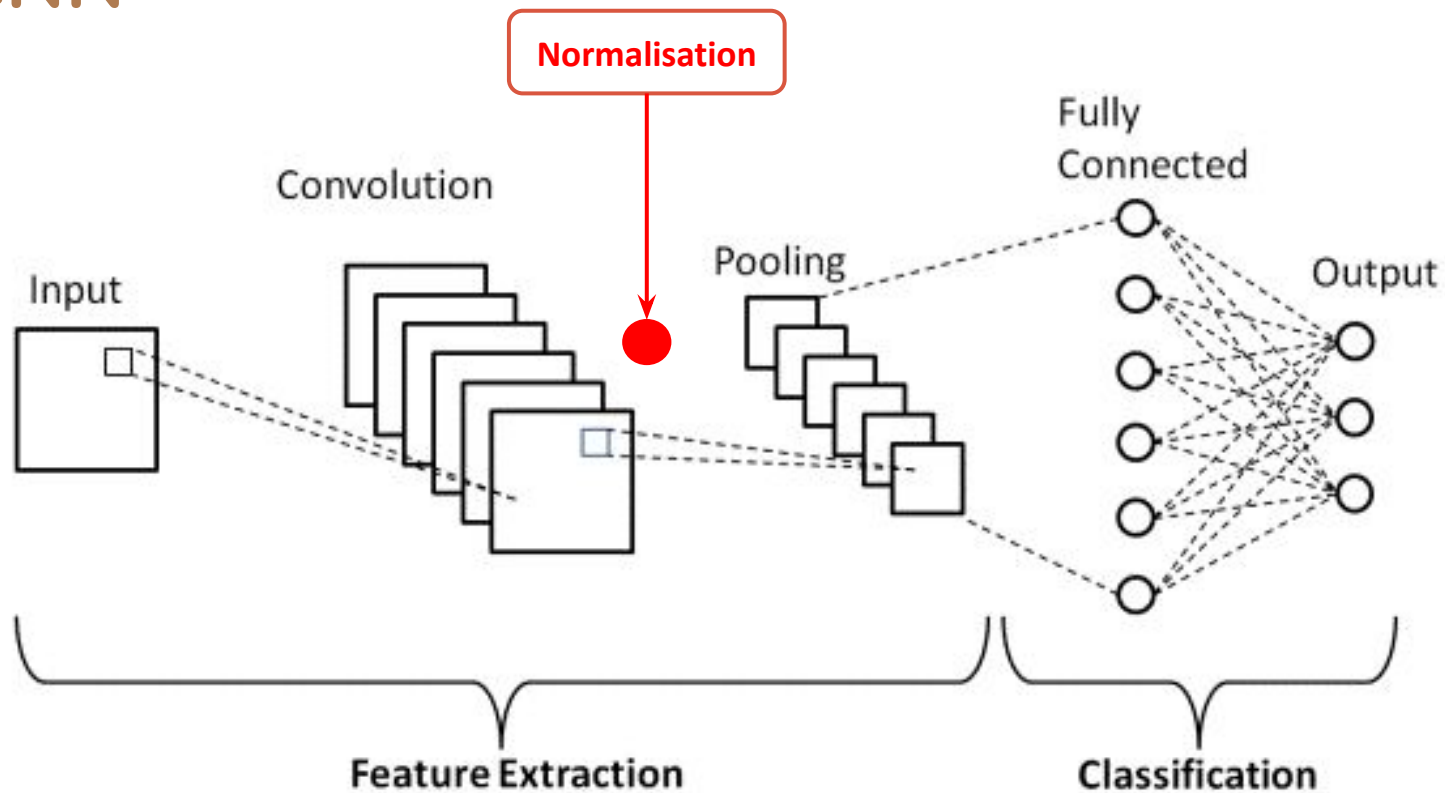


Figure 2. **Normalization methods.** Each subplot shows a feature map tensor, with  $N$  as the batch axis,  $C$  as the channel axis, and  $(H, W)$  as the spatial axes. The pixels in blue are normalized by the same mean and variance, computed by aggregating the values of these pixels.

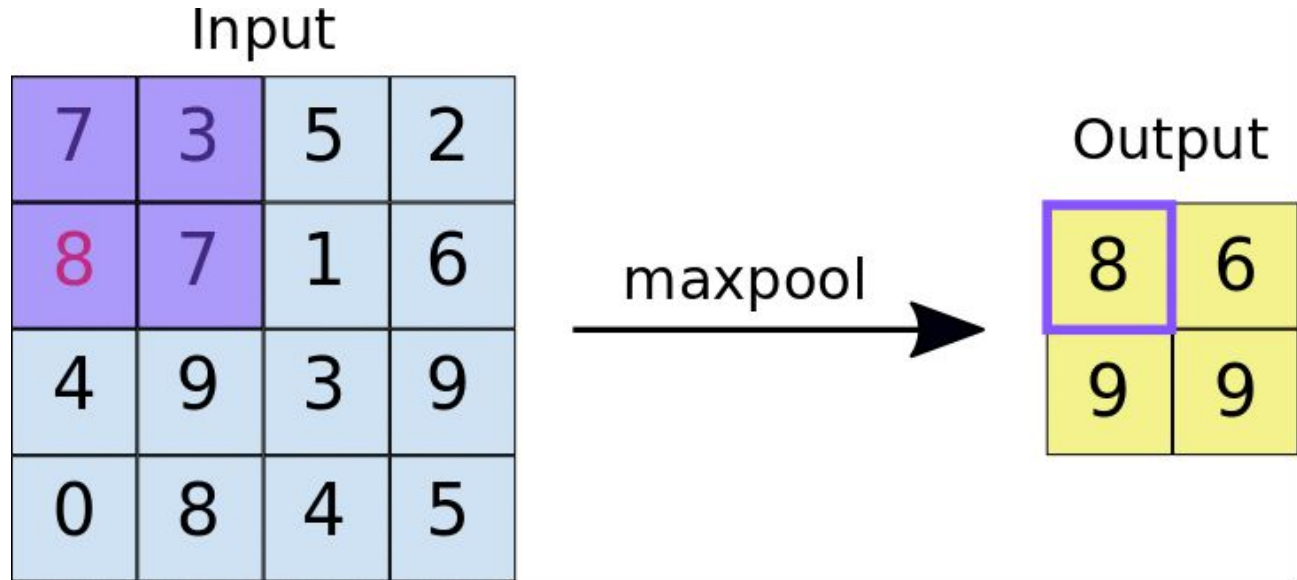
## Advantages:

- Improved Training Speed
- Better Gradient Flow: smooth the flow
- Reduced Sensitivity to Initialization
- Handling Different Batch Sizes

# CNN



# CNN: pooling layer



## Advantages:

- Dimension reduction
- Less sensitive to feature position

# CNN: pooling layer

## Feature maps (input):

- $W$  : width (pixels)
- $H$  : height (pixels)
- $D$  : previously used number of filters

## Hyperparameters:

- $F$  : cell size ( $F \times F$  pixels)
- $S$  : pixels between cells

## Feature maps (output):



- $W_p : (W - F) / S + 1$
- $H_p : (H - F) / S + 1$
- $D_p : D$

## Common choice:

- $F = 2$  or  $3$
- $S = 2$

# CNN: Features extraction

**Pooling layer output**

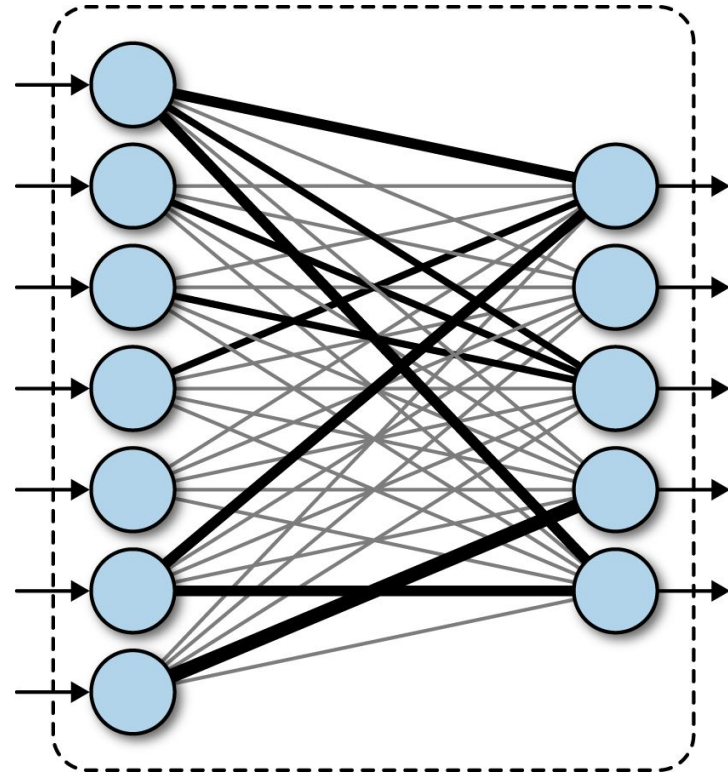




# CNN: fully-connected layer

## Options:

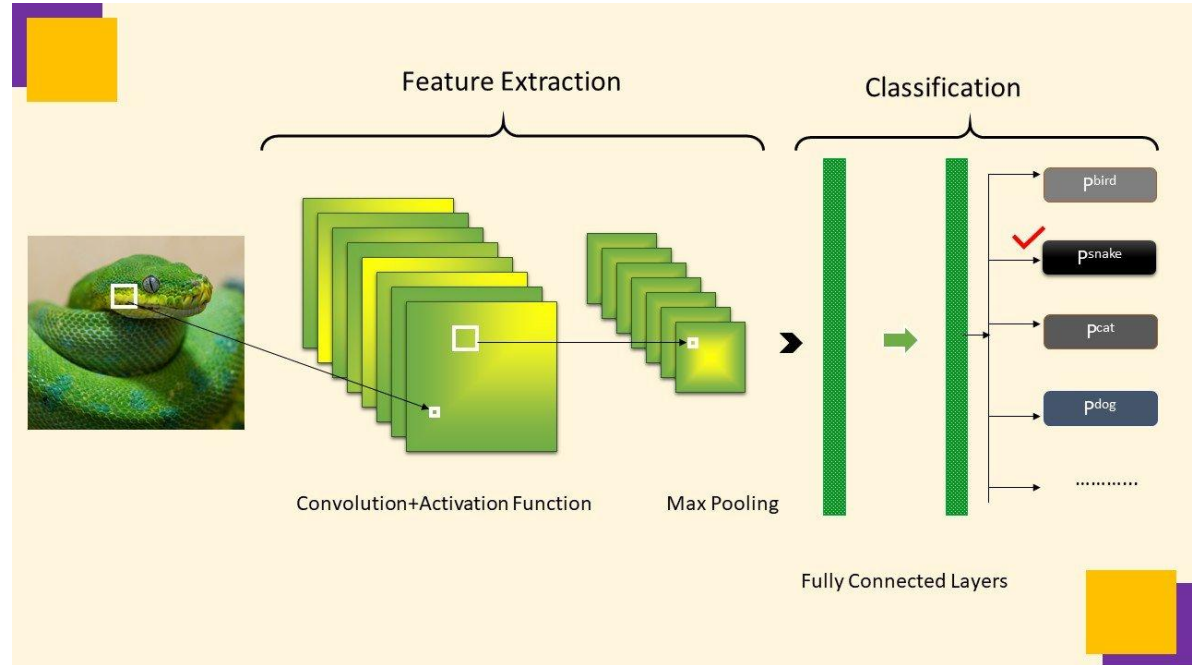
- Flatten: convert output to 1D vector
- Number of neurons
- Dropout: rate
- Activation:
  - ReLU
  - Classifier 2N: sigmoid
  - Classifier xN: Softmax




# CNN: summary

## Key points:

- Conv layers => extract features
- Pooling layers => reduce dimensions
- NN => classify



# CNN: VGG16



# CNN: VGG16

```
1  from keras.models import Sequential
2  from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
3
4  # Empty NN
5  my_VGG16 = Sequential()
6
7  #####
8  ### Convolution block ###
9  #####
10
11 # Add the first CNN layer with relu activation
12 my_VGG16.add(Conv2D(64, (3, 3), input_shape=(224, 224, 3), padding='same', activation='relu'))
13
14 # Add the second CNN layer with relu activation
15 my_VGG16.add(Conv2D(64, (3, 3), padding='same', activation='relu'))
16
17 # Add the first pooling layer
18 my_VGG16.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))
19
20 # Repeat the previous steps as often as needed...
21
22 #####
23 ### Fully-connected block ###
24 #####
25
26 # Convert 3D matrix to 1D vector
27 my_VGG16.add(Flatten())
28
29 # Add first fully-connected layer with relu activation
30 my_VGG16.add(Dense(4096, activation='relu'))
31
32 # Add second fully-connected layer with relu activation
33 my_VGG16.add(Dense(4096, activation='relu'))
34
35 # Add last fully-connected layer which is the classifier
36 my_VGG16.add(Dense(1000, activation='softmax'))
37
```

# Compilation

The suspense is intense !

# Hyperparameters

- Loss function
- Optimizer:
  - Learning rate
  - Momentum
- Metrics

```
67 #####
68 #####
69 ### Transfer learning ###
70 #####
71
72 from keras import Model
73
74 # Load the VGG16 trained on imagenet images without the fully-connected layer
75 model = VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3))
76
77 # Get the output of VGG16
78 x = model.output
79
80 # Add your new classifier
81 predictions = Dense(10, activation='softmax')(x)
82
83 # Design your new model
84 new_model = Model(inputs=model.input, outputs=predictions)
85
86 # Strat n°1:
87 for layer in new_model.layers:
88     layer.trainable = True
89
90 # Strat n°2:
91 for layer in new_model.layers:
92     layer.trainable = False
93
94 # Strat n°3: let's not train the first 5 layers (the largest)
95 for layer in new_model.layers[:5]:
96     layer.trainable = False
97
98 # Compile the new model
99 new_model.compile(loss="categorical_crossentropy", optimizer=optimizers.SGD(lr=0.0001, momentum=0.9), metrics=["accuracy"])
100
101 # Train it
102 model_info = new_model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size, verbose=2)
103
104
```

# Which loss ?

## ➤ **Classification:**

- Cross-entropy: Categorical / sparse

## ➤ **Object detection:**

- Intersection over Union (IoU) Loss
- Smooth L1 Loss => outliers robust

## ➤ **Generative Models:**

- Adversarial Loss: match the generated output distribution to the target distribution.
- Pixel-wise Mean Squared Error (MSE) Loss

## ➤ **Semantic Segmentation:**

- Pixel-wise Cross-Entropy Loss
- Dice Loss

## ➤ **Instance Segmentation:**

- Mask-RCNN Loss:
  - Cross-entropy
  - smooth L1
  - binary cross-entropy or Dice loss.

# Which optimizer ?

## ➤ **SGD:**

- Small dataset
- Small CNN architecture
- No LR scheduler
- Set LR and Momentum carefully

## ➤ **Adam and RMSprop:**

- LR scheduler
- LR and Momentum settings have little impact
- RMSprop slower but more stable
- RMSprop requires less RAM

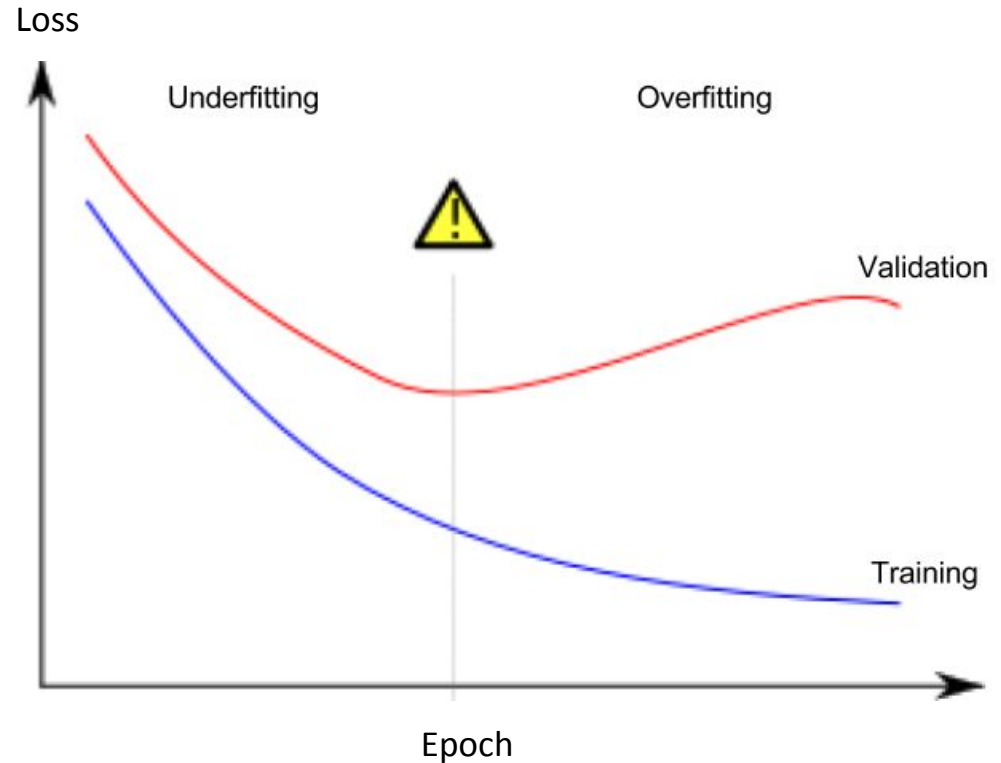


**Empirical testing  
required !**



# Underfitting vs. Overfitting

```
my_model.fit(X_train, y_train,  
             early_stopping_rounds=5,  
             eval_set=[(X_valid, y_valid)],  
             verbose=False)
```



# Transfer learning

Time to be lazy and... smart ! :)

# Transfer learning: purpose

## Weights:

- Convolution layer:
  - $F$  : filter size
  - $K$  : number of filters
  - $\text{weights} = F \times F \times K + K$
- Fully-connected:
  - $W$  : width of input
  - $H$  : height of input
  - $K$  : number of filters
  - $N$  : number of neurons
  - $\text{Weights} = W \times H \times K \times N + N$

**VGG16 = 138 357 544  
weights !!!**

# Transfer learning: strategy

## The idea:

- Keep the architecture
- Replace the classifier

## Strategy:

- Pre-trained model initialization
- Freezing initial layers
- Make the upper layers trainable
- Optional: adjusting upper layers: add / remove
- Training on the New Task

# Transfer learning: exemple

```
38 #####
39 #####
40 ### Pre-trained exemple ###
41 #####
42
43 from keras.applications.vgg16 import VGG16
44
45 # Get VGG16 available in keras
46 model = VGG16()
47
48 from keras.preprocessing.image import load_img, img_to_array
49 from keras.applications.vgg16 import preprocess_input
50
51 # Load image and reshape image to match VGG16 dimensions expectation
52 img = load_img('cat.jpg', target_size=(224, 224))
53
54 # Convert to numpy array because keras processes images as such
55 img = img_to_array(img)
56
57 # A CNN expects (in general) a collection of images.
58 # Reshape the array to add the number of images
59 img = img.reshape((1, img.shape[0], img.shape[1], img.shape[2]))
60
61 # Preprocess images the same way the images used to train VGG16 were preprocessed.
62 img = preprocess_input(img)
63
64 # Predict the class (or label) of this image
65 y = model.predict(img)
66
```

# Transfer learning: exemple

```
67 #####
68 ### Transfer learning ###
69 #####
70
71
72 from keras import Model
73
74 # Load the VGG16 trained on imagenet images without the fully-connected layer
75 model = VGG16(weights="imagenet", include_top=False, input_shape=(224, 224, 3))
76
77 # Get the output of VGG16
78 x = model.output
79
80 # Add your new classifier
81 predictions = Dense(10, activation='softmax')(x)
82
83 # Design your new model
84 new_model = Model(inputs=model.input, outputs=predictions)
85
86 # Strat n°1:
87 for layer in new_model.layers:
88     layer.trainable = True
89
90 # Strat n°2:
91 for layer in new_model.layers:
92     layer.trainable = False
93
94 # Strat n°3: let's not train the first 5 layers (the largest)
95 for layer in new_model.layers[:5]:
96     layer.trainable = False
97
98 # Compile the new model
99 new_model.compile(loss="categorical_crossentropy", optimizer=optimizers.SGD(lr=0.0001, momentum=0.9), metrics=["accuracy"])
100
101 # Train it
102 model_info = new_model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size, verbose=2)
103
104
```

# Transfer learning: tools



TensorFlow Hub

PYTORCH  
HUB



**Hugging Face**

# Examples


Architectures




# ResNet

## Residual connections:

- Faster training
- Better performances



# EfficiencyNet



EfficientNet Architecture

