

Reinforcement Learning III

Model-Free Control and Model-Based RL Approach

Antoine SYLVAIN

EPITA


2022-2023

Contents

- 1 Control in Monte-Carlo Methods
- 2 Control in Temporal-Difference Learning
- 3 Q-Learning

Policy Iteration (reminder)


Policy Iteration



Model-Free Policy Iteration using Action-Values

- Greedy policy improvement over $V(s)$ requires a model of the MDP
 $\pi' = \operatorname{argmax}_{a \in \mathcal{A}} \mathcal{R}_s^a + \mathcal{P}_{ss'}^a V(s')$
- Greedy policy improvement over $Q(s, a)$ is model-free:
 $\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a)$

Policy Iteration with action-value function



Policy Iteration with action-value function

Monte Carlo Generalized Policy Iteration


$$\pi_0 \rightarrow \pi_1 \rightarrow \pi_2 \rightarrow \dots$$

Improvement:

$$\pi_{k+1}(s) \doteq \operatorname{argmax}_a q_{\pi_k}(s, a)$$

Evaluation:

Monte Carlo Prediction



Monte-Carlo ES Algorithm

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize:

$\pi(s) \in \mathcal{A}(s)$ (arbitrarily), for all $s \in \mathcal{S}$

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Loop forever (for each episode):

Choose $S_0 \in \mathcal{S}$, $A_0 \in \mathcal{A}(S_0)$ randomly such that all pairs have probability > 0

Generate an episode from S_0, A_0 , following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$\pi(S_t) \leftarrow \arg\max_a Q(S_t, a)$

Example of blackjack



Example of blackjack



Example of blackjack

S,A	Returns(S,A)	Q(S,A)		$\bar{\pi}(S)$
		Hit	Stick	



Example of blackjack

S,A	Returns(S,A)	Q(S,A)		$\bar{\pi}(S)$
		Hit	Stick	
X,Stick	Returns(X,Stick) = [1]	0	1	Stick

$X = (\text{NoAce}, 20, 8)$



Example of blackjack



S,A	Returns(S,A)	Q(S,A)		$\pi(S)$
		Hit	Stick	
X,Stick	Returns(X,Stick) = [1]	0	1	Stick
Y,Hit	Returns(Y,Hit) = [1]	1	0	

$X = (\text{NoAce}, 20, 8)$

$Y = (\text{NoAce}, 13, 8)$



Example of blackjack



Limits of Exploring-Starts strategies

- We can not always chose the start
- We don't always to chose a new start
- How can we ensure exploration ?

Monte-Carlo ϵ -soft Algorithm

On-policy first-visit MC control (for ϵ -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small $\epsilon > 0$

Initialize:

$\pi \leftarrow$ an arbitrary ϵ -soft policy

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$


$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \operatorname{argmax}_a Q(S_t, a)$ (with ties broken arbitrarily)

For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \epsilon + \epsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \epsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

MC Control



Every episode:

Policy evaluation Monte-Carlo policy evaluation, $Q \approx q_\pi$

Policy improvement ϵ -greedy policy improvement

Contents

- 1 Control in Monte-Carlo Methods
- 2 Control in Temporal-Difference Learning
- 3 Q-Learning

Control with TD

- Natural idea: use TD instead of MC in our control loop
- Apply TD to $Q(S, A)$
- Use ϵ -greedy policy improvement
- Update every time-step


SARSA Algorithm

- SARSA
- $S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}$
- The agent chooses an action in the initial state to create the first state-action pair
- It observes the reward and next state
- Then it needs to know the next state-action pair to update its value estimates
- Since our agent is learning action values for a specific policy, it uses that policy to sample the next action:
$$S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1} \sim \pi$$

SARSA Algorithm

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

SARSA Improvement



Every **time-step**:

Policy evaluation **Sarsa**, $Q \approx q_\pi$

Policy improvement ϵ -greedy policy improvement

SARSA Algorithm

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'


 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Windy Gridworld Example




standard
moves



king's
moves

- Reward = -1 per time-step until reaching goal
- Undiscounted

Windy Gridworld Example



Contents

- 1 Control in Monte-Carlo Methods
- 2 Control in Temporal-Difference Learning
- 3 Q-Learning

On-Policy, Off-Policy

- **On-Policy:** improve and evaluate the policy being used to select actions
- **Off-Policy:** improve and evaluate a different policy from the one used to select actions

Off-Policy

- Learning Off-Policy, we use 2 policies:
 - **The target policy**: the policy we are learning, usually noted $\pi(a|s)$
 - **The behavior policy**: the one we use to select actions, usually noted $\mu(a|s)$
- Learn from observing other agents
- Re-use experience generated from old policies $\pi_1, \pi_2, \dots, \pi_{t-1}$
- Learn about optimal policy while following exploratory policy
- Learn about multiple policies while following one policy

Importance Sampling

- Sample : $x \sim \mu$
- Estimate : $\mathbb{E}_{\pi}[x]$

Derivation of Importance Sampling

$$\begin{aligned}\mathbb{E}_{\pi}[X] &\doteq \sum_{x \in X} x \pi(x) \\ \mathbb{E}_{\pi}[X] &= \sum_{x \in X} x \pi(x) \frac{\mu(x)}{\mu(x)} \\ \mathbb{E}_{\pi}[X] &= \sum_{x \in X} x \frac{\pi(x)}{\mu(x)} \mu(x) \\ \mathbb{E}_{\pi}[X] &= \sum_{x \in X} x \rho(x) \mu(x)\end{aligned}$$

$\rho(x) = \frac{\pi(x)}{\mu(x)}$ is called the importance sampling ratio

Derivation of Importance Sampling

$$\mathbb{E}_\pi[X] = \sum_{x \in X} x \rho(x) \mu(x)$$

$$\mathbb{E}_\pi[X] = \mathbb{E}_\mu[X \rho(X)]$$

$$\mathbb{E}_\mu[X \rho(X)] = \sum_{x \in X} x \rho(x) \mu(x)$$

Using a weighted sample average $\mathbb{E} \approx \frac{1}{n} \sum_{i=1}^n x_i$

$$\mathbb{E}_\mu[X \rho(X)] \approx \sum_{i=1}^n x_i \rho(x_i) \approx \mathbb{E}_\pi[X]$$

$$x \sim \mu$$

Importance Sampling for Off-Policy MC

- Use returns generated from μ to evaluate π
- Weight return G_t according to similarity between policies
- Multiply importance sampling corrections along whole episode

$$G_t^{\pi/\mu} = \frac{\pi(A_t, S_t)}{\mu(A_t, S_t)} \frac{\pi(A_{t+1}, S_{t+1})}{\mu(A_{t+1}, S_{t+1})} \cdots \frac{\pi(A_\tau, S_\tau)}{\mu(A_\tau, S_\tau)}$$

- Update value towards corrected return
$$V(S_t) \leftarrow V(S_t) + \alpha(G_t^{\pi/\mu} - V(S_t))$$
- Cannot use if μ is zero when π is non-zero
- Importance sampling can dramatically increase variance


Importance Sampling for Off-Policy TD

- Use TD targets generated from μ to evaluate π
- Weight TD target $R + \gamma V(S')$ by importance sampling
- Only need a single importance sampling correction
$$V(S_t) \leftarrow V(S_t) + \alpha \left(\frac{\pi(A_t, S_t)}{\mu(A_t, S_t)} (R_{t+1} + \gamma V(S_{t+1})) - V(S_t) \right)$$
- Much lower variance than Monte-Carlo importance sampling
- Policies only need to be similar over a single step

Q-Learning

- We now consider off-policy learning of action-values $Q(s, a)$
- No importance sampling is required
- Next action is chosen using behaviour policy $A_{t+1} \sim \mu(.|S_t)$
- But we consider alternative successor action $A' \sim \pi(.|S_t)$
- And update $Q(S_t, A_t)$ towards value of alternative action
$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$$

Q-Learning Control



$$Q(S, A) \leftarrow Q(S, A) + \alpha \left(R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right)$$

Q-Learning Algorithm

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

Q-Learning Algorithm

Initialized

Q-Table		Actions					
		South (0)	North (1)	East (2)	West (3)	Pickup (4)	Dropoff (5)
States	0	0	0	0	0	0	0

	327	0	0	0	0	0	0

	499	0	0	0	0	0	0

Training



Q-Table		Actions					
		South (0)	North (1)	East (2)	West (3)	Pickup (4)	Dropoff (5)
States	0	0	0	0	0	0	0

	328	-2.30108105	-1.97092096	-2.30357004	-2.20591839	-10.3607344	-8.5583017

	499	9.96984239	4.02706992	12.96022777	29	3.32877873	3.38230603


Q-Learning in Windy Gridworld Example

The Windy Gridworld





A specific example: the Cliff environment

The cliff walking environment



A specific example: the Cliff environment

The cliff walking environment



Thank you for your attention